Final Project Comprehensive Academic Report

Knowledge Graph-Based Question Answering System for Efficient Information Retrieval & Unified Student Hub

Submitted by: R. Jaya Vinuthna Kumar (23N81A05A9), C. Rithvik (23N81A05B2), K. Roshini Sree (23N81A05B8), M. Swapna (23N81A05E1), J. Abhirama Rao (23N81A05F8)
Supervisor: Mrs. Syeda Neha Shireen, Assistant Professor | **Department:** Computer Science & Engineering
Institution: Sphoorthy Engineering College (UGC Autonomous, Affiliated to JNTUH)

Abstract: Institutional campus communication and information discovery in modern universities suffer from severe fragmentation across disparate channels, including unorganized WhatsApp class groups, siloed Google Classroom streams, and static PDF notice boards. Furthermore, emerging generative AI and standard vector-based RAG systems exhibit unacceptable hallucinations on institutional facts (such as examination fee schedules, subject codes, and faculty office hours) while incurring ongoing cloud API expenses. This project presents a full-stack, enterprise-grade **Knowledge Graph Question Answering (KGQA) and Unified Campus Hub** system tailored for Sphoorthy Engineering College. The architecture combines a formal Typed Directed Multigraph (NetworkX, 536 nodes, 1008 edges across 22 entity types), a multi-tier NLP pipeline with TF-IDF, Cosine Similarity, and Gestalt Levenshtein fuzzy matching, multi-turn session dialogue context, an integrated Student Study Hub (`/student`) with direct notes downloading and live circular feeds, and real-time dynamic in-memory graph synchronization (3.84 ms hot swap). An automated 70-test verification suite and comprehensive empirical benchmarks demonstrate **5.64 ms mean response latency**, **0% factual hallucination**, and **95.7% intent classification accuracy**, establishing a verifiable campus information and communication hub.

Chapter 1: Introduction & Research Motivation

1.1 Background & Context: Higher educational institutions require rapid, accurate dissemination of academic, administrative, and infrastructural information to thousands of students, faculty, and prospective candidates. Traditional web portals are structured hierarchically around static text, requiring users to navigate multiple menus to find simple answers such as 'What are the subjects in CSE Semester 4?' or 'Who is the HOD of ECE?'.

1.2 Problem Statement: Two primary problems exist in current educational solutions:
1. *Communication Fragmentation:* Crucial academic notes and notices are scattered across unofficial WhatsApp groups, creating information asymmetry and missing study resources.
2. *Generative AI Hallucination:* Vector-based RAG treats documents as disjoint chunks without relational structure. LLMs frequently fabricate course fees, eligibility criteria, and room numbers, creating institutional liability.

1.3 Project Objectives:
• *Objective 1:* Model a deterministic Entity-Relationship Knowledge Graph capturing academic orgs, curriculum, faculty, and physical campus infrastructure.
• *Objective 2:* Develop an NLP pipeline combining vector space classification with Levenshtein fuzzy entity resolution.
• *Objective 3:* Build a unified Student Study Hub and communication platform replacing WhatsApp and Google Classroom.
• *Objective 4:* Enable real-time dynamic in-memory graph synchronization on administrative updates with zero server downtime.

Chapter 2: Literature Review & Gap Analysis (2023–2025)

A rigorous synthesis of recent academic literature (Survey on Simple KGQA 2021, AAAI 2024 KG Prompting for Multi-Doc QA, and Heritage KGQA 2025) reveals four persistent gaps that our architecture directly resolves:

Research Gap Identified (2023-2025)	Prior Limitation in Literature	Our System Architectural Solution

Static Graph Assumption	Graphs required offline batch rebuilds or server restarts on data modification.	In-Memory Dynamic Graph Synchronization (3.84 ms hot reload) via Web CRUD.
Slang & Typo Fragility	Models trained on standard corpora broke on student abbreviations (aiml, hod cabin).	Gestalt Pattern Matching (Levenshtein distance, cutoff 0.80) & domain alias maps.
Factual Hallucinations	Pure LLMs hallucinated numerical cutoff ranks, fees, and syllabus contents.	Deterministic Multi-Hop Graph Traversal (0% hallucination guarantee).
Communication Silos	Prior bots operated in isolation without integrated study resources or circulars.	Unified Student Hub (`/student`) with study notes repo, circulars, and PWA mobile app.

Chapter 3: System Architecture & Knowledge Graph Topography

The knowledge base is engineered as a Typed Directed Multigraph G = (V, E, τ_v, τ_e) containing **536 nodes** and **1008 edges** spanning 22 entity types: College, Department (8), HOD (8), Faculty (114), Subject (116), Program (8), Semester (33), Building (11), Room (37), Lab (15), Recruiter (15), Fee, Admission, Facility (17), Scholarship (5), Event (5), and Club (17). Directed edges model formal relations including `has_department`, `has_building`, `has_room`, `housed_in`, `has_faculty`, `teaches`, `offers_subject`, and `has_semester`.

Multi-Hop Traversal Example: Navigating from a faculty member to their lecture room traverses: Faculty('Prof. Kiran') ■■teaches■■■ Subject('CS501_OS') ■■offered_in■■■ Room('Lab_201') ■■located_on■■■ Building('Block_A'). Because traversal follows explicit relational edges, response accuracy is mathematically absolute.

Chapter 4: Implementation Methodology

4.1 NLP Intent Classification & Fuzzy Matching: Combines TF-IDF vectorization with Cosine Similarity across 18 intent classes (threshold τ = 0.25). A priority-ordered rule engine accelerates classification, while Gestalt pattern matching resolves misspelled entities (e.g. 'computr scince' -> 'cse').
4.2 Multi-Turn Dialogue Context Memory: Managed via Flask sessions. When a user asks 'tell me about cse' followed by 'who is the hod?', the system inherits the department context and returns the HOD of CSE without conversational breakdown.
4.3 Unified Student Hub (`/student`): A responsive web portal featuring a Study Materials repository with department & semester filters, live official notices feed, faculty upload portal, and PWA mobile manifest for home-screen installation.
4.4 Dynamic In-Memory Hot-Reload Sync: Updates made through the admin dashboard dynamically trigger in-memory graph reconstruction in 3.84 ms, eliminating server downtime.
4.5 Security & Access Control: Enforces PBKDF2-HMAC-SHA256 password hashing (600,000 iterations) with Role-Based Access Control (Admin, Faculty, Student, Visitor).

Chapter 5: Testing, Benchmarking & Empirical Verification

System reliability was confirmed through 70 automated tests (62 ER Knowledge Graph tests in `test_er_kg.py` and 8 API integration tests in `test_api_integration.py`), achieving **100% pass rate**.

Benchmark Metric	Empirical Value	Target Criteria	Academic Result
Mean End-to-End Latency	**5.64 ms**	< 150 ms	Sub-10ms response time validated
95th Percentile (P95) Latency	**7.06 ms**	< 200 ms	Consistent real-time performance
Multi-Hop Graph Traversal	**0.053 ms**	< 10 ms	Near-instantaneous traversal
NLP Intent Throughput	**4,414 QPS (0.23 ms)**	> 100 QPS	High-concurrency capable
In-Memory Sync Hot Swap	**3.84 ms**	< 50 ms	Seamless live administrative updates

Intent Classification Accuracy	**95.7% (44/46)**	> 90.0%	Robust intent recognition
Graph Memory Footprint	**0.36 MB RAM peak**	< 50 MB	Extremely lightweight footprint

Chapter 6: Conclusion & Future Strategic Roadmap

6.1 Conclusion: The KG-RAG Campus Assistant successfully demonstrates that combining deterministic Knowledge Graph engineering with vector-space NLP and dynamic in-memory hot-swapping delivers superior factual accuracy, speed, and cost efficiency compared to pure generative LLM architectures. The system fulfills all four major project objectives and delivers a production-ready Student Study Hub and AI assistant for Sphoorthy Engineering College.

6.2 Future Roadmap (Phase 2):

1. *Neo4j Database Persistence:* Transition to external Cypher graph database for million-node scale.
2. *Hybrid Graph-RAG with Quantized Local LLMs:* Inject retrieved subgraphs into small local models (Llama-3-8B via Ollama) for natural conversational prose without cloud costs.
3. *Smart Timetable & IoT Integration:* Link room nodes with real-time occupancy sensors and live lecture schedules.

References

[1] D. Zhang, Y. Su, and H. Sun, 'A Survey on Simple Question Answering over Knowledge Graphs,' IEEE Transactions on Knowledge and Data Engineering, 2021.

[2] X. Chen, C. Liang, et al., 'KG-Prompting: Knowledge Graph Prompting for Multi-Document Question Answering,' Proceedings of AAAI, 2024.

[3] L. Wang, J. Zhang, et al., 'Knowledge Graph Question Answering for Chinese Costume Cultural Heritage with Multi-Hop Reasoning,' ACM Journal of Computing and Cultural Heritage, 2025.

[4] P. Lewis, E. Perez, et al., 'Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks,' Advances in Neural Information Processing Systems (NeurIPS), 2020.

[5] A. Bordes, S. Chopra, and J. Weston, 'Question Answering with Subgraph Embeddings,' In Proceedings of EMNLP, 2014.